

Sketch demoProject highlights

Gustavo Homem

homem.gustavo@gmail.com October 2023

TLDR - migration tool

```
gustavo@sketch-operations:~/sketch-demo/migration$ python3 sketch migrate.py -b 100 -p 10
WARNING: this script will modify rows at database proddatabase after copying the files at sketch-legacy-s3 to sketch-production-s3.
Are you sure you want to continue? (yes/no) yes
Execution starting
Connecting to the database
Connecting to the S3 storage
Checking $3 read permissions for sketch-legacy-s3 on fra1.digitaloceanspaces.com
Checking S3 write permissions for sketch-production-s3 on fra1.digitaloceanspaces.com
Current data status:
 * 2001 objects in bucket sketch-legacy-s3
 * 1999 objects in bucket sketch-production-s3
 * 4000 total objects found
 * 2001 legacy entries in the database
 * 1999 production entries in the database
 * O unexpected entries in the database
 * 4000 total entries in the database
NOTE: this script does not delete legacy objects
Are you sure you want perform the migration over this data status? (yes/no) yes
Migrating legacy data
The log file for this execution will be /tmp/sketch_migrate_gustavo_2023-11-02-19-05_Q2kCVt.log
Progress information:
 * Progress 48%, batches processed 10/21, files copied 1000, rows updated 1000, elapsed time 15.89
 * Progress 95%, batches processed 20/21, files copied 2000, rows updated 2000, elapsed time 39.25
 * Progress 100%, batches processed 21/21, files copied 2001, rows updated 2001, elapsed time 39.58
Execution finished after 39.58 seconds
Current data status:
 * 2001 objects in bucket sketch-legacy-s3
 * 4000 objects in bucket sketch-production-s3
 * 6001 total objects found
 * 0 legacy entries in the database
 * 4000 production entries in the database
 * 0 unexpected entries in the database
 * 4000 total entries in the database
NOTE: this script does not delete legacy objects
The log file can be reviewed with:
less /tmp/sketch_migrate_gustavo_2023-11-02-19-05_Q2kCVt.log
gustavo@sketch-operations:~/sketch-demo/migration$
```

TLDR - migration tool

gustavo@sketch-operations:~/sketch-demo/migration\$ python3 sketch migrate.py -b 100 -p 10 Confirm you want WARNING: this script will modify rows at database proddatabase after copying the files at sketch-legacy-s3 to sketch-production-s3. Are you sure you want to continue? (yes/no) yes to write into prod Execution starting Connecting to the database Connecting to the S3 storage Checking \$3 read permissions for sketch-legacy-s3 on fra1.digitaloceanspaces.com Checking S3 write permissions for sketch-production-s3 on fra1.digitaloceanspaces.com Current data status: * 2001 objects in bucket sketch-legacy-s3 * 1999 objects in bucket sketch-production-s3 * 4000 total objects found * 2001 legacy entries in the database * 1999 production entries in the database * 0 unexpected entries in the database * 4000 total entries in the database NOTE: this script does not delete legacy objects Confirm the initial Are you sure you want perform the migration over this data status? (yes/no) yes situation makes Migrating legacy data The log file for this execution will be /tmp/sketch_migrate_gustavo_2023-11-02-19-05_Q2kCVt.log sense Progress information: * Progress 48%, batches processed 10/21, files copied 1000, rows updated 1000, elapsed time 15.89 Check progress * Progress 95%, batches processed 20/21, files copied 2000, rows updated 2000, elapsed time 39.25 * Progress 100%, batches processed 21/21, files copied 2001, rows updated 2001, elapsed time 39.58 Execution finished after 39.58 seconds Current data status: * 2001 objects in bucket sketch-legacy-s3 * 4000 objects in bucket sketch-production-s3 Review the final * 6001 total objects found situation * 0 legacy entries in the database * 4000 production entries in the database * 0 unexpected entries in the database * 4000 total entries in the database NOTE: this script does not delete legacy objects The log file can be reviewed with: Be aware of the log less /tmp/sketch_migrate_gustavo_2023-11-02-19-05_Q2kCVt.log gustavo@sketch-operations:~/sketch-demo/migration\$ file

TLDR - migration tool in verbose mode

```
* 488 production entries in the database
 * O unexpected entries in the database
 * 1000 total entries in the database
NOTE: this script does not delete legacy objects
Are you sure you want perform the migration over this data status? (yes/no) yes
Migrating legacy data
The log file for this execution will be /tmp/sketch migrate gustavo 2023-11-02-19-36 tzHWym.log
Progress information:
The execution of SELECT took 0.0 seconds
Got S3 batch
 * copying sketch-legacy-s3/image/avatar-000000980.png to sketch-production-s3/avatar/avatar-000000980.png
 * copying sketch-legacy-s3/image/avatar-000000981.png to sketch-production-s3/avatar/avatar-000000981.png
 * copying sketch-legacy-s3/image/avatar-000000983.png to sketch-production-s3/avatar/avatar-000000983.png
 * copying sketch-legacy-s3/image/avatar-000000984.png to sketch-production-s3/avatar/avatar-000000984.png
 * copying sketch-legacy-s3/image/avatar-000000986.png to sketch-production-s3/avatar/avatar-000000986.png
 * copying sketch-legacy-s3/image/avatar-000000987.png to sketch-production-s3/avatar/avatar-000000987.png
 * copying sketch-legacy-s3/image/avatar-000000989.png to sketch-production-s3/avatar/avatar-000000989.png
 * copying sketch-legacy-s3/image/avatar-000000991.png to sketch-production-s3/avatar/avatar-000000991.png
 * copying sketch-legacy-s3/image/avatar-000000992.png to sketch-production-s3/avatar/avatar-000000992.png
 * copying sketch-legacy-s3/image/avatar-000000993.png to sketch-production-s3/avatar/avatar-000000993.png
 * copying sketch-legacy-s3/image/avatar-000000995.png to sketch-production-s3/avatar/avatar-000000995.png
 * copying sketch-legacy-s3/image/avatar-000000997.png to sketch-production-s3/avatar/avatar-000000997.png
S3 batch done
The execution of copy s3 batch took 1.22 seconds, avg 0.1 per file
Got DB batch
 * updating image/avatar-000000980.png to avatar/avatar-000000980.png
 * updating image/avatar-000000981.png to avatar/avatar-000000981.png
 * updating image/avatar-000000983.png to avatar/avatar-000000983.png
 * updating image/avatar-000000984.png to avatar/avatar-000000984.png
 * updating image/avatar-000000986.png to avatar/avatar-000000986.png
 * updating image/avatar-000000987.png to avatar/avatar-000000987.png
 * updating image/avatar-000000989.png to avatar/avatar-000000989.png
 * updating image/avatar-000000991.png to avatar/avatar-000000991.png
 * updating image/avatar-0000000992.png to avatar/avatar-000000992.png
 * updating image/avatar-000000993.png to avatar/avatar-000000993.png
 * updating image/avatar-000000995.png to avatar/avatar-000000995.png
 * updating image/avatar-000000997.png to avatar/avatar-000000997.png
DB batch done
The execution of update db batch took 0.01 seconds, avg 0.0 seconds per row
```

The verbose information is available both on screen and in the log file.

(but in the log file lines are prefixed with timestamps)

TLDR – simulated environment preparation tool

```
ustavo@sketch-operations:~/sketch-demo/preparation$ python3 sketch prepare.py -c 4000
                                                                                                                                                                       Confirm you want
WARNING: this script will DROP and re-CREATE database proddatabase and delete all contents of bucket sketch-legacy-s3 and bucket sketch-production-s3
Are you sure you want to continue? (yes/no) yes
                                                                                                                                                                       reset the simulation
Execution starting
Re-creating the database
                                                                                                                                                                       data
Initializing the database
 * user migration created with password Av6muiP9GuSl for migration purposes
Cleaning the legacy bucket
 * partial count: found 533 objects, deleted 533 objects
                                                                                                                                                                       Check progress
 * final count: found 533 objects, deleted 533 objects
Cleaning the production bucket
 * partial count: found 1000 objects, deleted 1000 objects
 * final count: found 1000 objects, deleted 1000 objects
Creating S3 objects and the corresponding database rows
                                                                                                                                                                       Review the final
Created 2001 legacy avatars and 1999 production avatars
 ustavo@sketch-operations:~/sketch-demo/preparation$
                                                                                                                                                                       situation
```

TLDR – integration test



Compliance

Please check the compliance.md document for a one by one review of the requirements.

Agenda

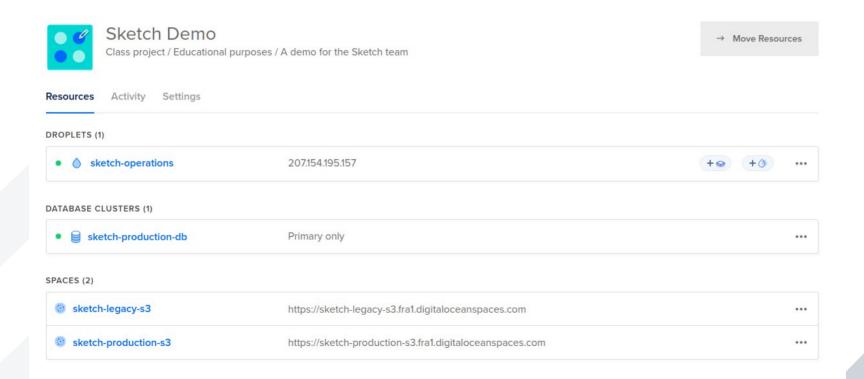
- Environment
- Deployment
- Performance

Environment

The execution environment is a Digital Ocean project comprised of:

- An Ubuntu Linux machine
- A PostgreSQL database cluster
- Two storage Spaces (S3 compatible buckets)

Environment



Environment

I chose Digital Ocean because:

- the montly Spaces subscription includes 1TB of bandwidth (performance testing generates many requests and AWS charges per request)
- performance results would not be realistic on local S3 emulation environments

Side effect: an opportunity to test the S3 API compatibility of a different provider.

Agenda

- Environment
- Deployment
- Performance

The deployment of the migration script is extremely simple because we only need an Ubuntu machine and the guarantee that 3 apt packages from the Ubuntu distribution are present.

```
node 'sketch-operations' {
    include is_puppet_base::node_base
    include passwd_common

$sketch_pkgs = [ 'postgresql-client', 'python3-boto3', 'python3-psycopg2' ]
    package { $sketch_pkgs : ensure => present, provider => apt }
}
```

```
node 'sketch-operations' {
    include is_puppet_base::node_base
    include passwd_common

$sketch_pkgs = [ 'postgresql-client', 'python3-boto3', 'python3-psycopg2' ]
    package { $sketch_pkgs : ensure => present, provider => apt }
}
```

The deployment was tested in the smallest Digital Ocean configuration (1 vCPU, 512M RAM) and nothing bigger than that is necessary.

The script could simply be deployed by hand using git clone or it could be deployed via IaC code. I can provide details on IaC deployment, if necessary.

Agenda

- Environment
- Deployment
- Performance

- Execute on a Linux instance near the database
- Avoid transversing the entire database more than once
- Do not overwrite files that are already migrated

- Execute on a Linux instance near the database
- Avoid transversing the entire database more than once
- Do not overwrite files that are already migrated

"your computer" just from this, for single process executions.

- Execute on a Linux instance near the database
- Avoid transversing the entire database more than once
- Do not overwrite files that are already migrated

3x improvement versus
 "your computer" just from
 this, for single process
 executions.

The problem is very latency-sensitive because we are copying a large number of small files. For large files it would be different.

- Execute on a Linux instance near the database
- Avoid transversing the entire database more than once
- Do not overwrite files that are already migrated

It is not only client to API latency that counts. With small files there is also the TCP slow start inside the provider.

 3x improvement versus "your computer" just from this, for single process executions.

The problem is very latency-sensitive because we are copying a large number of small files. For large files it would be different.

- Execute on a Linux instance near the database
- Avoid transversing the entire database more than once
- Do not overwrite files that are already migrated

Theoretically important and adopted in the implementation but it turns out that the SELECTs with the LIKE operator are very fast compared to the S3 copy operations.

- Execute on a Linux instance near the database
- Avoid transversing the entire database more than once
- Do not overwrite files that are already migrated

From ~0.07 seconds per file to less than 0.005 seconds (more than 14x faster).

Note: some files might be already migrated if we are re-executing over a the result of previous partial migration

All the low hanging fruit has been collected before the implementation of more complex performance optimizations. A simple implementation with all the low hanging fruit included delivered a modest migration performance of ~ 14.5 entries per second.

With this implementation it was possible to conclude that:

- time spent on the single used SELECT is negligible
- time spent on UPDATEs is small compared to the total time
- time spent on S3 copies dominates the process

With this implementation it was possible to conclude that:

- time spent on the single used SELECT is negligible
- time spent on UPDATEs is small compared to the total time
- time spent on S3 copies dominates the process

We already lowered the latency by using a dedicated instance that lives near the database.

With this implementation it was possible to conclude that:

- time spent on the single used SELECT is negligible
- time spent on UPDATEs is small compared to the total time
- time spent on S3 copies dominates the process

The S3 API does not implement bulk copies in a simple way (2)

With this implementation it was possible to conclude that:

- time spent on the single used SELECT is negligible
- time spent on UPDATEs is small compared to the total time
- time spent on S3 copies dominates the process

I chose to parallelize the copies with the standard python multiprocessing module.

- scan the database and pick X batches of Y entries to migrate
 - distribute X batches to X processes and wait for completion
 - pick the execution results from the processes using a shared queue
 - log progress and stats

- scan the database and pick X batches of Y entries to migrate
 - distribute X batches to X processes and wait for completion
 - pick the execution results from the processes using a shared queue
 - log progress and stats

The waiting part might seem strange but we need to ensure that we don't lose control of launched processes and we need to control the logging order of the migration steps. Furthermore, these processes will take roughly the same time to finish, therefore the waiting time of finished processes will be small compared to their execution time.

- scan the database and pick X batches of Y entries to migrate
 - distribute X batches to X processes and wait for completion
 - pick the execution results from the processes using a shared queue
 - log progress and stats
- implement an integration test that allows for detailed performance testing at low cognitive effort and find optimal values for X and Y

I executed the migration of 5000 entries 64 times on order to collect a modest set of performance data. If the process was not fully automated this would not have been possible within the given timeframe.

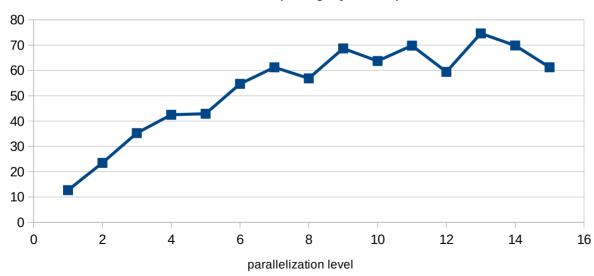
- scan the database and pick X batches of Y entries to migrate
 - distribute X batches to X processes and wait for completion
 - pick the execution results from the processes using a shared queue
 - log progress and stats
- implement an integration test that allows for detailed performance testing at low cognitive effort and find optimal values for X and Y

X and Y are exposed to the user as CLI parameters. See:

```
python3 sketch_migrate.py -h
```

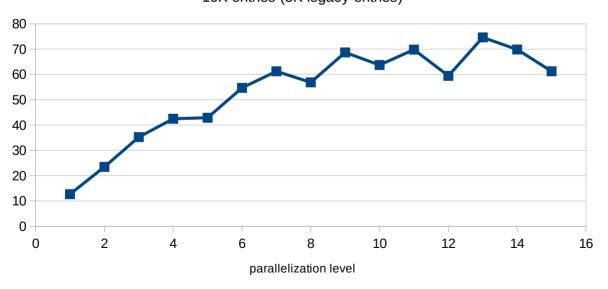
Performance vs parallelization level - batch size 100

10K entries (5K legacy entries)



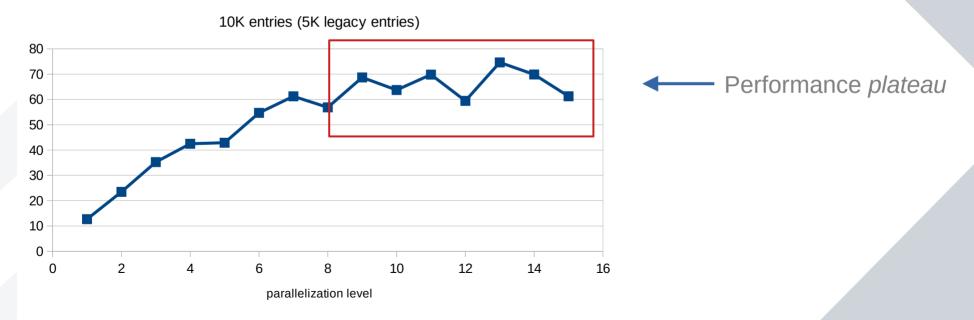
Performance vs parallelization level - batch size 100

10K entries (5K legacy entries)



 As usual, real world data is noisy, but the pattern is clear.

Performance vs parallelization level - batch size 100



- 4x 5x performance improvement
- the estimated performance is ~ 67 entries per second
- optimal batch size found to lie between 100 and 1000
- "optimal" parallelization level is between 5 and 10
 - there is a dramatic improvement as we increase the parallelization level from small numbers but a plateau is hit between 5 and 10

At 67 entries per second we would expect to migrate 1M entries in 4.15 h.

Tradeoffs



- No tradeoffs for the parallelization level: use at least of 10 processes as they fit perfectly even on a very small machine
- Small batches are less efficient because processes and network connections are started / stopped more often but large batches amplify the inter-batch waiting time for synchronization
- Small batches provide more frequent on screen progress updates than large batches
- If in doubt, use a batch size of 100

Hypothetical other steps for performance

If this performance is not good enough we could still consider:

- dissecting further where time is being consumed in the S3 copy processes
- investigating the complex S3 Batch operations, as an alternative
- check if individual Linux instances are rate-limited in the S3 API and, if so, split the process across multiple instances

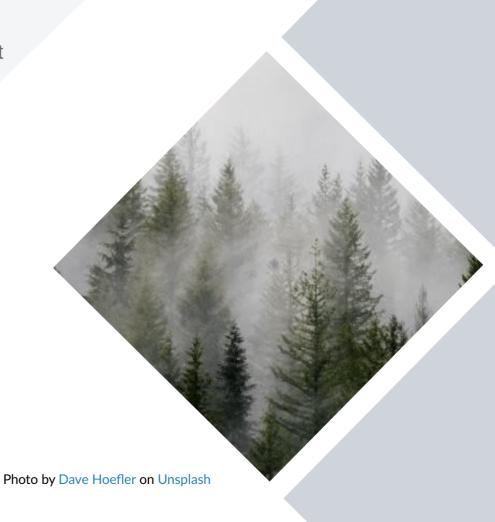
Each additional performance gain will come at the cost of additional R&D time and additional code complexity.

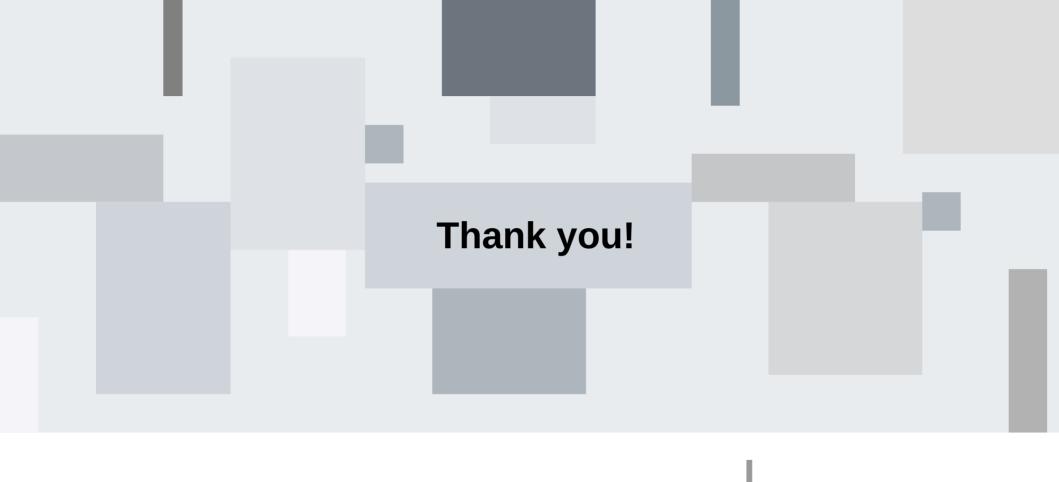
Final note

The performance results are orientative at this point. That is, these results allow us to understand the order of magnitude where we are working.

For the results to be statistically more robust we should perform multiple runs for every combination of batch size and parallelization level and calculate the average and standard deviation. There is intrinsinc variance in the process due to the variable load on the provider API.

The effort to invest in further performance work depends on how many millions of rows have to be migrated and how time sensitive is the process, given that there is no business discountinuity as the database entries are only updated after sucessful copy of the corresponding file to the new bucket.





Sketch demoProject highlights

Gustavo Homem

homem.gustavo@gmail.com October 2023